



Malignant Comments Classifier Project

Submitted by:
Raghavulu Patnala

ACKNOWLEDGMENT

Thanks for giving me the opportunity to work in Flip n Robo Technologies as Intern and would like to express my gratitude to Data Trained Institute as well for trained me in Data Science Domain.

This helps me to do my projects well and understand the concepts.

INTRODUCTION

- Business Problem Framing

The proliferation of social media enables people to express their opinions widely online. However, at the same time, this has resulted in the emergence of conflict and hate, making online environments uninviting for users. Although researchers have found that hate is a problem across multiple platforms, there is a lack of models for online hate detection.

Online hate, described as abusive language, aggression, cyberbullying, hatefulness and many others has been identified as a major threat on online social media platforms. Social media platforms are the most prominent grounds for such toxic behaviour.

There has been a remarkable increase in the cases of cyberbullying and trolls on various social media platforms. Many celebrities and influences are facing backlashes from people and must come across hateful and offensive comments. This can take a toll on anyone and affect them mentally leading to depression, mental illness, self-hatred and suicidal thoughts.

Internet comments are bastions of hatred and vitriol. While online anonymity has provided a new outlet for aggression and hate speech, machine learning can be used to fight it. The problem we sought to solve was the tagging of internet comments that are aggressive towards other users. This means that insults to third parties such as celebrities will be tagged as unoffensive, but “u are an idiot” is clearly offensive.

Our goal is to build a prototype of online hate and abuse comment classifier which can be used to classify hate and offensive comments so that it can be controlled and restricted from spreading hatred and cyberbullying.

- **Conceptual Background of the Domain Problem**

In social media the people spreading or involved in such kind of activities uses filthy languages, aggression, images etc. to offend and gravely hurt the person on the other side.

This is one of the major concerns now. The result of such activities can be dangerous. It gives mental trauma to the victims making their lives miserable.

Online hate, described as abusive language, aggression, cyberbullying, hatefulness, insults, personal attacks, provocation, racism, sexism, threats, or toxicity has been identified as a major threat on online social media platforms.

These kinds of activities must be checked for a better future.

- **Motivation for the Problem Undertaken**

The project was the first provided to me by Flip-Robo as a part of the internship programme. The exposure to real world data and the opportunity to deploy my skillset in solving a real time problem has been the primary objective.

The main aim is to build a prototype of online hate and abuse comment classifier which can be used to classify hate and offensive comments so that it can be controlled and restricted from spreading hatred and cyberbullying.

Analytical Problem Framing

- Mathematical/ Analytical Modeling of the Problem

Here we need to find whether the given comments are malignant words or not. It is text classification problem where we need to predict the target variable from the text and, we have multiple target variables like malignant, high malignant, rude, abuse, loathe.

- Data Sources and their formats

The Data is provided by Flip Robo Technologies, and it has Train and Test Data Set and need to train our data in Train dataset and need to load the Test dataset to make the predictions.

```
# Reading and Loading the train dataset.  
df_train=pd.read_csv('train.csv')  
df_train
```

	id	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0
...
159566	ffe987279560d7ff	"::::And for the second time of asking, when ...	0	0	0	0	0	0
159567	ffea4adeee384e90	You should be ashamed of yourself \n\nThat is ...	0	0	0	0	0	0
159568	ffe36eab5c267c9	Spitzer \n\nUmm, theres no actual article for ...	0	0	0	0	0	0
159569	fff125370e4aaaf3	And it looks like it was actually you who put ...	0	0	0	0	0	0
159570	fff46fc426af1f9a	"\nAnd ... I really don't think you understand...	0	0	0	0	0	0

```
# Reading and Loading the test dataset.  
df_test=pd.read_csv('test.csv')  
df_test
```

	id	comment_text
0	00001cee341fdb12	Yo bitch Ja Rule is more succesful then you'll...
1	0000247867823ef7	== From RfC == \n\n The title is fine as it is...
2	00013b17ad220c46	" \n\n == Sources == \n\n * Zawe Ashton on Lap...
3	00017563c3f7919a	:If you have a look back at the source, the in...
4	00017695ad8997eb	I don't anonymously edit articles at all.
...
153159	ffcd0960ee309b5	. \n i totally agree, this stuff is nothing bu...
153160	fffd7a9a6eb32c16	== Throw from out field to home plate. == \n\n...
153161	ffda9e8d6fafa9e	" \n\n == Okinotorishima categories == \n\n I ...
153162	ffe8f1340a79fc2	" \n\n == ""One of the founding nations of the...
153163	ffffce3fb183ee80	" \n ::Stop already. Your bullshit is not wel...

153164 rows × 3 columns

- Data Pre-processing Done

For Data pre-processing we did some data cleaning, where we used WordNet lemmatizer to clean the words and removed special characters using Regexp Tokenizer.

Then, filtered the words by removing stop words and then used lemmatizers and joined and return the filtered words.

Used TFIDF vectorizer to convert those text into vectors and trained the train and loaded the test dataset.

```
#Defining the stop words
stop_words = stopwords.words('english')

#Defining the Lemmatizer
lemmatizer = WordNetLemmatizer()

#Replacing '\n' in comment_text
df_train['comment_text'] = df_train['comment_text'].replace('\n', ' ')

#Function Definition for using regex operations and other text preprocessing for getting c
def clean_comments(text):

    #convert to lower case
    lowered_text = text.lower()

    #Replacing email addresses with 'emailaddress'
    text = re.sub(r'^.+@[^\.\.]*\.[a-z]{2,}$', 'emailaddress', lowered_text)

    #Replace URLs with 'webaddress'
    text = re.sub(r'http\S+', 'webaddress', text)

    #Removing numbers
    text = re.sub(r'[0-9]', " ", text)

    #Removing the HTML tags
    text = re.sub(r"<.*?>", " ", text)

    #Removing Punctuations
    text = re.sub(r'[\W\s]', ' ', text)
    text = re.sub(r'\_', ' ', text)

    #Removing all the non-ascii characters
    clean_words = re.sub(r'^\x00-\x7f$', r'', text)

    #Removing the unwanted white spaces
    text = " ".join(text.split())

    #Splitting data into words
    tokenized_text = word_tokenize(text)

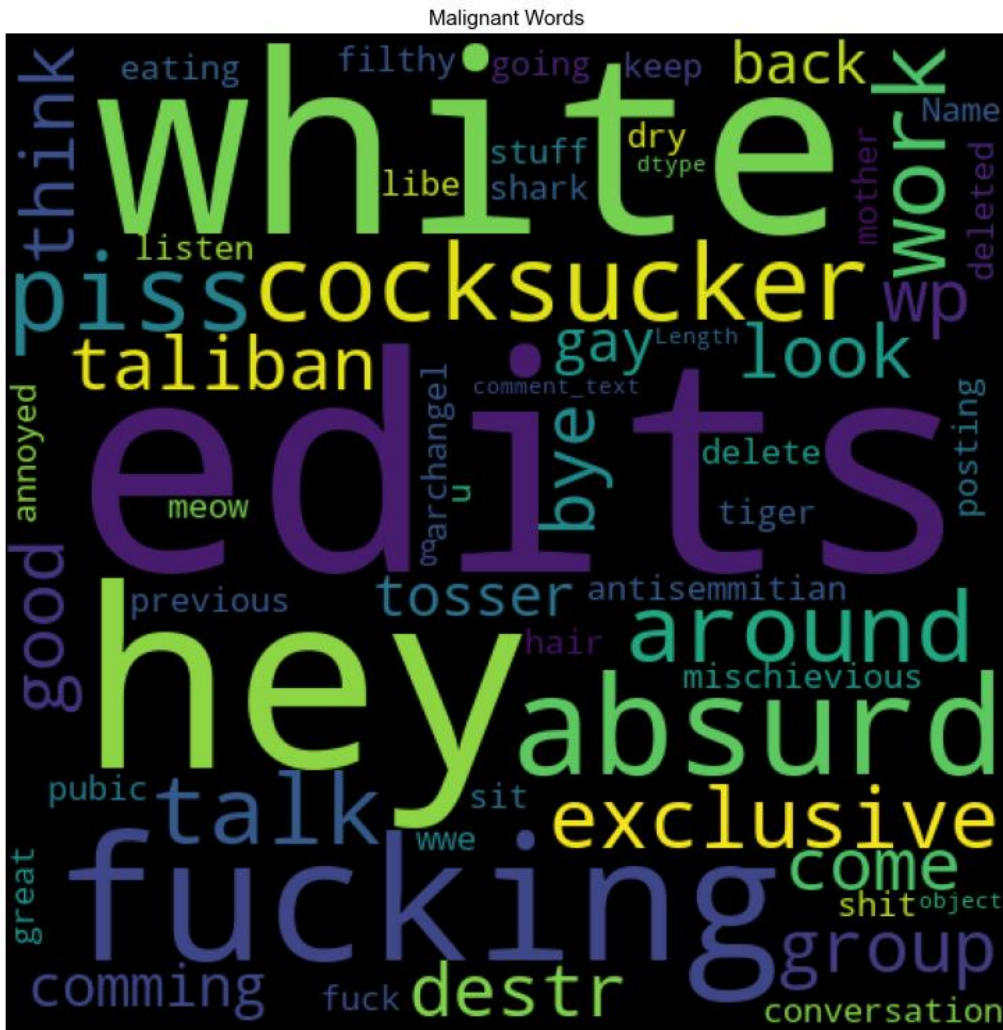
    #Removing remaining tokens that are not alphabetic, Removing stop words and Lemmatizin
    removed_stop_text = [lemmatizer.lemmatize(word) for word in tokenized_text if word not

    return " ".join(removed_stop_text)
```

```
#Converting the features into number vectors
tf_vec = TfidfVectorizer(max_features = 10000, stop_words='english')
```

```
#Let's Separate the input and output variables represented by X and y respectively in train
X = tf_vec.fit_transform(df_train['comment_text'])
```

- Data Inputs- Logic- Output Relationships



A word cloud visualization of words found in the comments section of a video titled "How to make a cock sucking sound". The words are arranged in a dense, overlapping manner, with colors ranging from dark purple to bright green. The most prominent words include "shit", "white", "stuff", "work", "stupid", "sideways", "fucking", "edit", "absurd", "cock sucker", "dirty", "gay", "asshole", "fist fucker", "around", "piss", "meow", "antisemitian", "love", "posting", "mother", "dry", "libe", "gg", "comment_text", "shark", "keep length", "conversation", "sit", "great", "tony", "arm", "peace", "deleted", "hair", "wwe", "archangel", "hey", "object", "Name", "public", "mischievous", "listen", "stop", "fuck", "previous", "annoyed", "tiger", "going", "eating", "deleting", "obviously", "surrounding", "fist fucker", "absurd", "cock sucker", "dirty", "gay", "asshole", "fist fucker", "absurd", "cock sucker", "dirty", "gay", "asshole".

[illegible]

From the above graph we can see the most used words in all categories – malignant, highly malignant, abuse, loathe, rude.

- Hardware and Software Requirements and Tools Used

Model training was done on Jupiter Notebook. Kernel Version is Python3.

Hardware -- > Intel 8GB RAM, i5 processor

```
# Importing all the required libraries.

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from collections import Counter
import string
import re

# packages from gensim
from gensim import corpora
from gensim.parsing.preprocessing import STOPWORDS
from gensim.utils import simple_preprocess

# packages from sklearn
from sklearn.feature_extraction.text import TfidfVectorizer

# packages from nltk
import nltk
from nltk.corpus import wordnet
from nltk.corpus import stopwords
from wordcloud import WordCloud
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer, SnowballStemmer
from nltk import pos_tag

import warnings
warnings.filterwarnings('ignore')
```

The above libraries and packages used in this project for building a model.

```
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score, RandomizedSearchCV
from sklearn.metrics import f1_score, accuracy_score, classification_report, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import PassiveAggressiveClassifier
```

Model/s Development and Evaluation

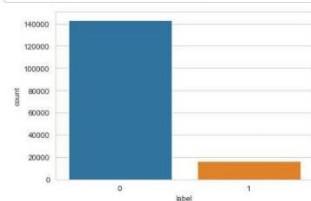
- Identification of possible problem-solving approaches (methods)

Converting the label into 0 and 1 as below,

```
# Creating a new feature having Negative Comments and Non-Negative Comments from all features
df_train['label'] = df_train[features].max(axis=1)
df_train.head(10)
```

	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe	label
0	Explanation\nWhy the edits made under my userm...	0	0	0	0	0	0	0
1	D'aww! He matches this background colour I'm S...	0	0	0	0	0	0	0
2	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0	0
3	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0	0
4	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0	0
5	"\n\nCongratulations from me as well, use the ...	0	0	0	0	0	0	0
6	COCKSUCKER BEFORE YOU PISS AROUND ON MY WORK	1	1	1	0	1	0	1
7	Your vandalism to the Matt Shrivington article...	0	0	0	0	0	0	0
8	Sorry if the word 'nonsense' was offensive to ...	0	0	0	0	0	0	0
9	alignment on this subject and which are contra...	0	0	0	0	0	0	0

```
plt.figure(figsize = (6,4))
sns.countplot(df_train['label'])
plt.show()
df_train['label'].value_counts()
```



```
0    143346
1     16225
Name: label, dtype: int64
```

- Testing of Identified Approaches (Algorithms)
 - Logistic Regression
 - Gradient Boost Classifier
 - Decision Tree Classifier
 - Naïve Bayes Multi-Nomial NB
 - Passive Aggressive Classifier

- Run and evaluate selected models

```

lor = LogisticRegression()
lor.fit(x_train,y_train)
y_pred = lor.predict(x_test)
scr_lor = cross_val_score(lor,x_over,y_over,cv=5)

print("F1 score \n", f1_score(y_test,y_pred))
print("CV Score :", scr_lor.mean())
print("-----\n")
print("Classification Report \n", classification_report(y_test,y_pred))
print("-----\n")
print("Confusion Matrix \n", confusion_matrix(y_test,y_pred))
print("ROC AUC Score \n", roc_auc_score(y_test,y_pred))

```

```

F1 score
0.9320920043811611
CV Score : 0.9313340260855669
-----

Classification Report
precision    recall  f1-score   support

     0       0.94      0.92      0.93      35600
     1       0.92      0.94      0.93      36073

 accuracy          0.93      0.93      0.93      71673
 macro avg          0.93      0.93      0.93      71673
 weighted avg          0.93      0.93      0.93      71673

-----

Confusion Matrix
[[32673  2927]
 [ 2033 34040]]
ROC AUC Score
0.9307114790171116

```

```

gb = GradientBoostingClassifier()
gb.fit(x_train,y_train)
y_pred = gb.predict(x_test)
scr_gb = cross_val_score(gb,x_over,y_over,cv=5)

print("F1 score \n", f1_score(y_test,y_pred))
print("CV Score :", scr_gb.mean())
print("-----\n")
print("Classification Report \n", classification_report(y_test,y_pred))
print("-----\n")
print("Confusion Matrix \n", confusion_matrix(y_test,y_pred))
print("ROC AUC Score \n", roc_auc_score(y_test,y_pred))

```

```

F1 score
0.8067577098159902
CV Score : 0.8335391739704592
-----

Classification Report
precision    recall  f1-score   support

     0       0.76      0.97      0.85      35600
     1       0.96      0.70      0.81      36073

 accuracy          0.83      0.83      0.83      71673
 macro avg          0.86      0.83      0.83      71673
 weighted avg          0.86      0.83      0.83      71673

-----

Confusion Matrix
[[34451  1149]
 [10907 25166]]
ROC AUC Score
0.8326828069766146

```

```

dt = DecisionTreeClassifier()
dt.fit(x_train,y_train)
y_pred = dt.predict(x_test)
scr_dt = cross_val_score(dt,x_over,y_over,cv=5)

print("F1 score \n", f1_score(y_test,y_pred))
print("CV Score :", scr_dt.mean())
print("-----\n")
print("Classification Report \n", classification_report(y_test,y_pred))
print("-----\n")
print("Confusion Matrix \n", confusion_matrix(y_test,y_pred))
print("ROC AUC Score \n", roc_auc_score(y_test,y_pred))

```

```

F1 score
0.9472996337794235
CV Score : 0.9483487853559177
-----

```

```

Classification Report
precision    recall  f1-score   support

     0       0.96      0.93      0.94     35600
     1       0.93      0.96      0.95     36073

 accuracy          0.95     71673
 macro avg         0.95     71673
weighted avg         0.95     71673
-----

```

```

Confusion Matrix
[[33011  2589]
 [ 1282 34791]]
ROC AUC Score
0.9458681175375651

```

```

mnmb= MultinomialNB()
mnmb.fit(x_train,y_train)
y_pred = mnmb.predict(x_test)
scr_mnmb = cross_val_score(mnmb,x_over,y_over,cv=5)

print("F1 score \n", f1_score(y_test,y_pred))
print("CV Score :", scr_mnmb.mean())
print("-----\n")
print("Classification Report \n", classification_report(y_test,y_pred))
print("-----\n")
print("Confusion Matrix \n", confusion_matrix(y_test,y_pred))
print("ROC AUC Score \n", roc_auc_score(y_test,y_pred))

```

```

F1 score
0.8938483307415345
CV Score : 0.8978276387083366
-----

```

```

Classification Report
precision    recall  f1-score   support

     0       0.89      0.90      0.89     35600
     1       0.90      0.88      0.89     36073

 accuracy          0.89     71673
 macro avg         0.89     71673
weighted avg         0.89     71673
-----

```

```

Confusion Matrix
[[32195  3405]
 [ 4172 31901]]
ROC AUC Score
0.894349782525883

```

```

pac = PassiveAggressiveClassifier()
pac.fit(x_train,y_train)
y_pred = pac.predict(x_test)
scr_pac = cross_val_score(pac,x_over,y_over,cv=5)

print("F1 score \n", f1_score(y_test,y_pred))
print("CV Score :", scr_pac.mean())
print("-----\n")
print("Classification Report \n", classification_report(y_test,y_pred))
print("-----\n")
print("Confusion Matrix \n", confusion_matrix(y_test,y_pred))
print("ROC AUC Score \n", roc_auc_score(y_test,y_pred))

```

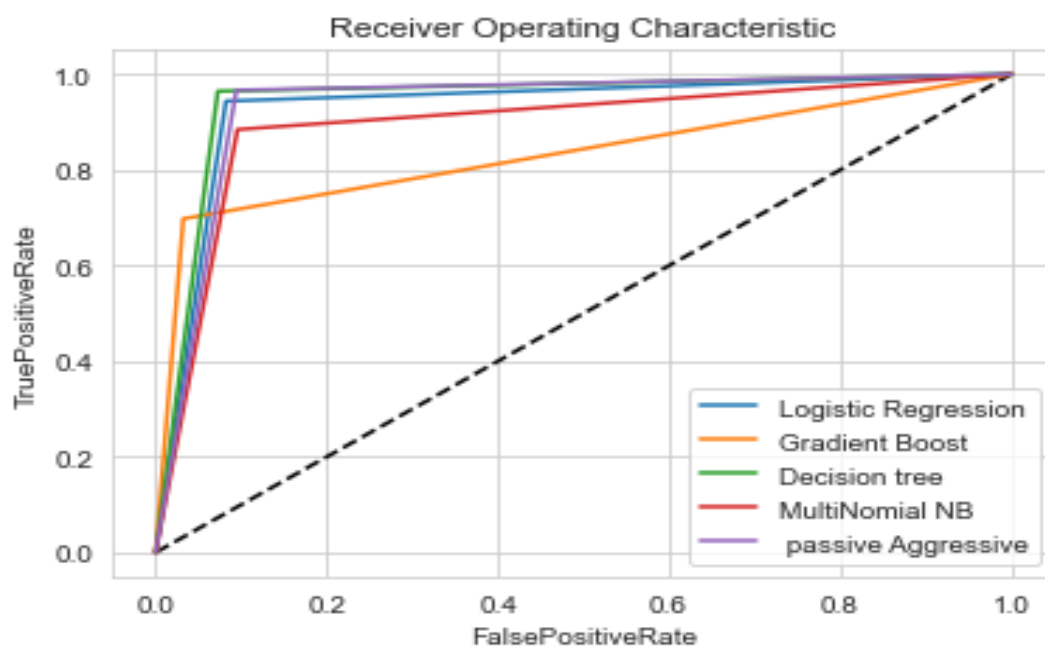
F1 score
 0.9385892897998734
 CV Score : 0.9374869473254114

Classification Report				
	precision	recall	f1-score	support
0	0.96	0.91	0.93	35600
1	0.91	0.97	0.94	36073
accuracy			0.94	71673
macro avg	0.94	0.94	0.94	71673
weighted avg	0.94	0.94	0.94	71673

Confusion Matrix
 [[32240 3360]
 [1203 34870]]
 ROC AUC Score
 0.9361344676540735

- Key Metrics for success in solving problem under consideration

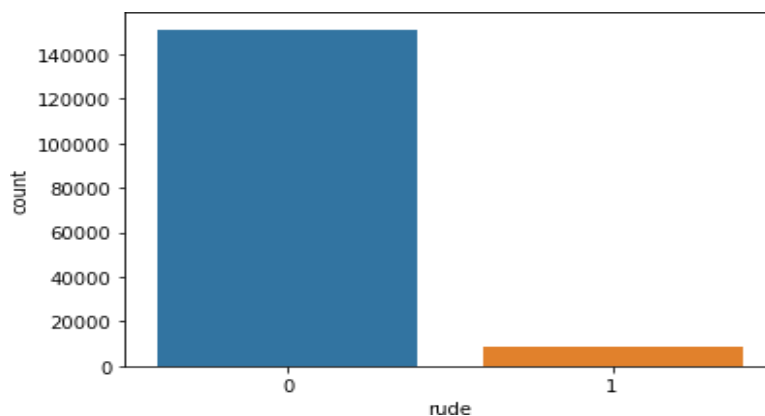
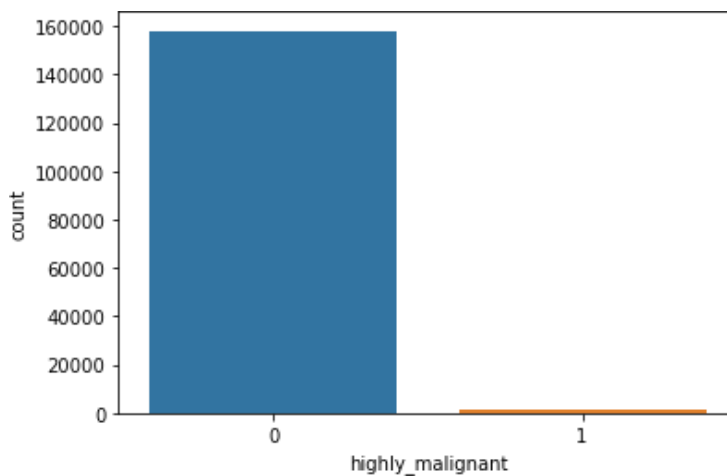
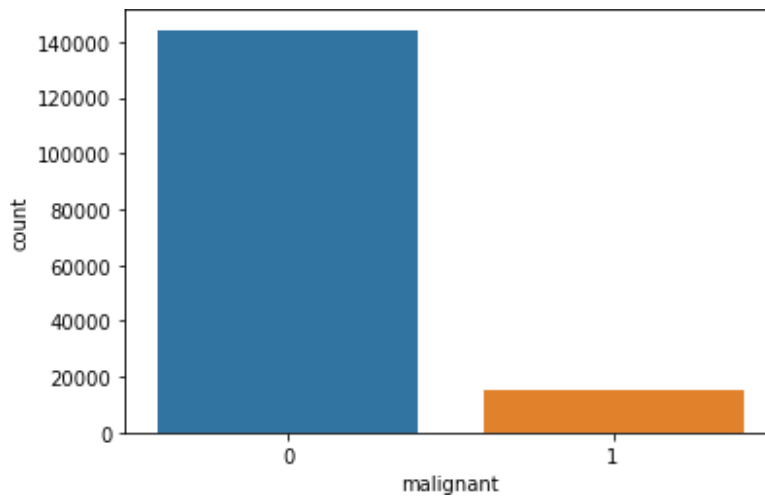
Key Metrics used were the Accuracy Score, Cross validation Score and AUC & ROC Curve as this was binary classification as you can see in the above image in models used.

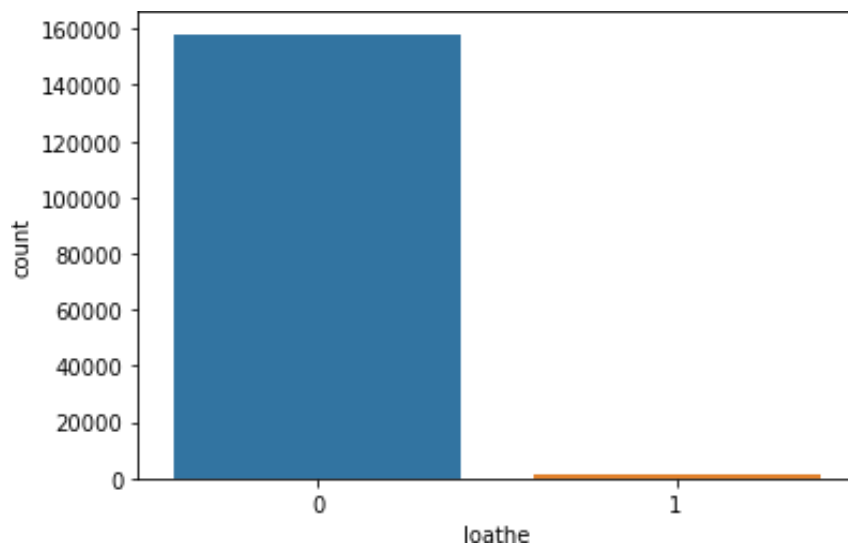
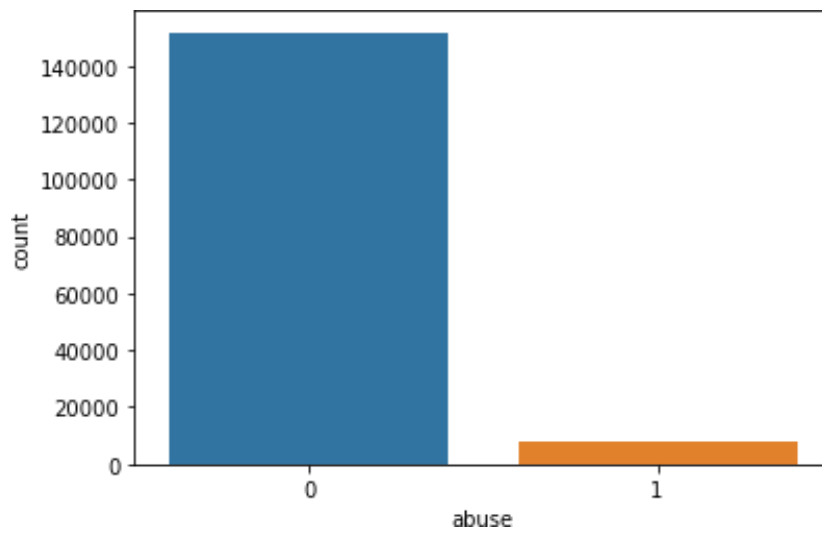
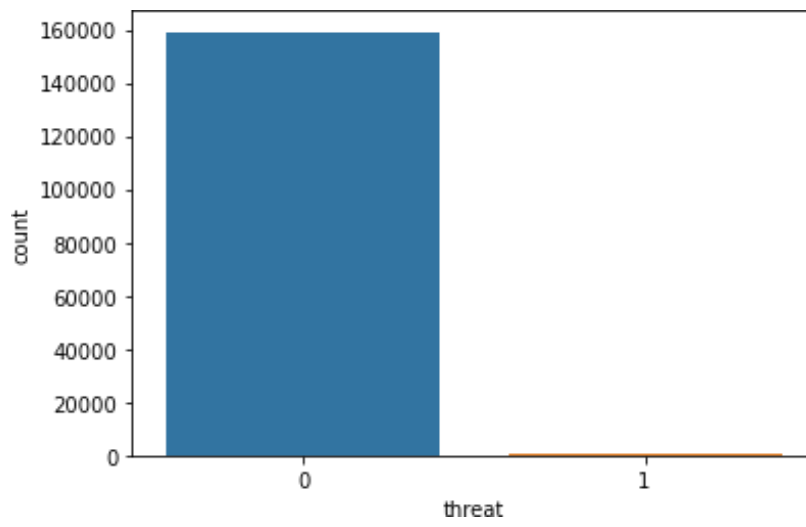


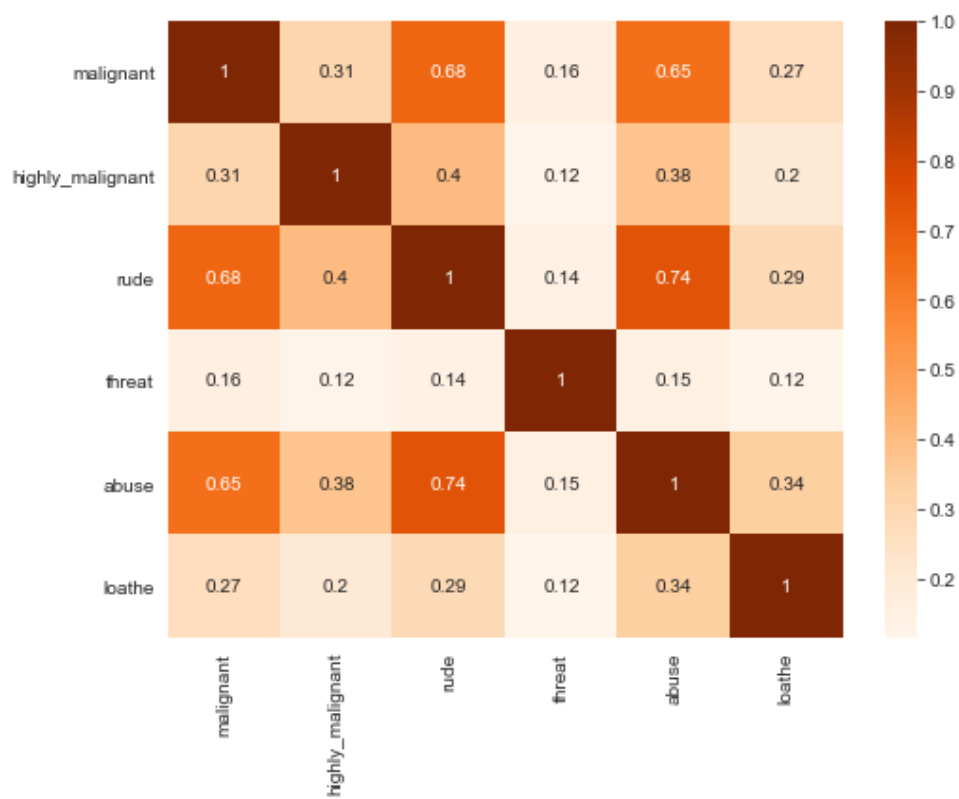
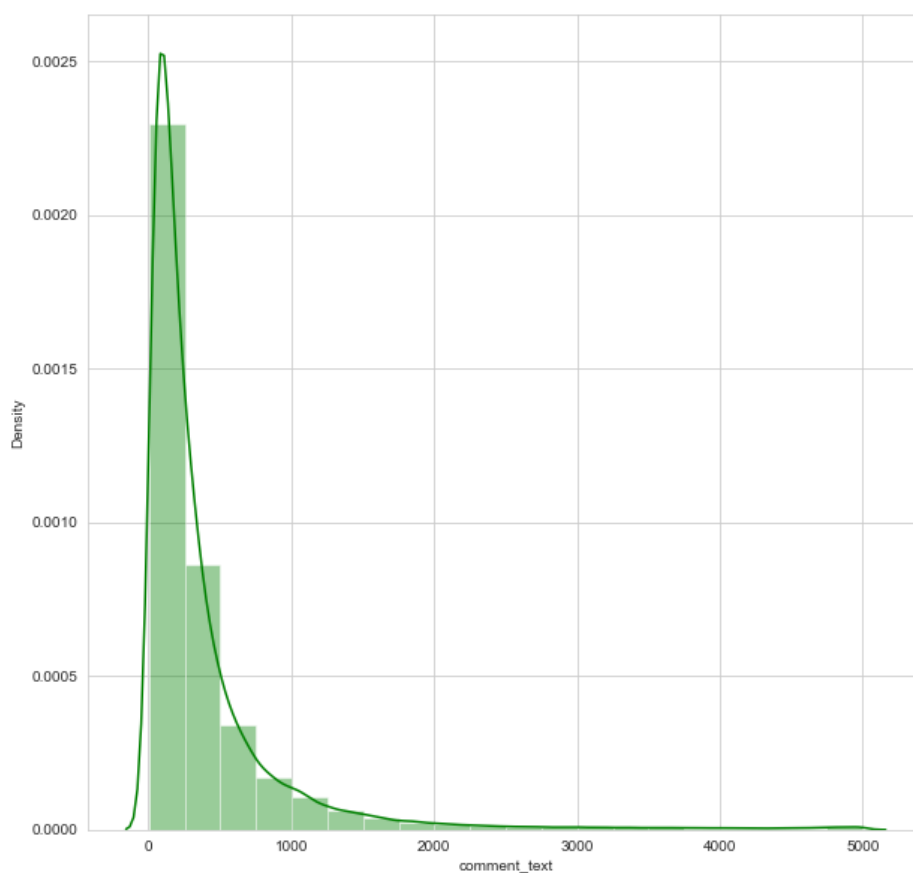
- Visualizations

Used Count plot and distribution plot and for the different target variables.

Heat map for test the correlation between features and variables.







- Interpretation of the Results

```
# Lets try to improve the accuracy of model by hyper parameter tuning,

param = {'C': [1.0,1.2,1.4,1.6,1.8],
         'fit_intercept':[True], 'max_iter': [1000]}

# Applying randomized search CV to increase the accuracy,

rg = RandomizedSearchCV(pac, param_distributions = param, cv= 5)
rg.fit(x_train,y_train)
rg.best_params_

{'max_iter': 1000, 'fit_intercept': True, 'C': 1.0}
```

```
#final model accuracy,

model = PassiveAggressiveClassifier(C = 1.0, max_iter = 1000, fit_intercept = True)

model.fit(x_train,y_train)
y_pred = model.predict(x_test)

print("F1 score \n", f1_score(y_test,y_pred))
print("-----\n")
print("Classification Report \n", classification_report(y_test,y_pred))
print("-----\n")
print("Confusion Matrix \n", confusion_matrix(y_test,y_pred))
print("ROC AUC Score \n", roc_auc_score(y_test,y_pred))
```

```
F1 score
0.9374109423309585
-----

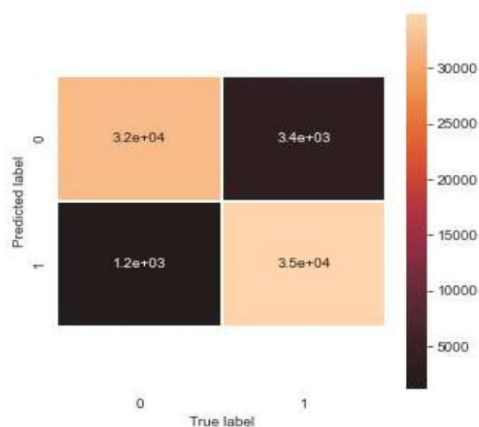
Classification Report
              precision    recall  f1-score   support

         0       0.96       0.90       0.93       35600
         1       0.91       0.97       0.94       36073

 accuracy       0.94
 macro avg       0.94       0.93       0.93       71673
 weighted avg    0.94       0.94       0.93       71673
-----

Confusion Matrix
[[32150  3450]
 [ 1206 34867]]
ROC AUC Score
0.9348288403633456
```

```
# Confusion matrix Visualization
fig, ax = plt.subplots(figsize=(5,5))
sns.heatmap(confusion_matrix(y_test, y_pred),annot=True,linewidths=1,center=0)
plt.xlabel("True label")
plt.ylabel("Predicted label")
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
```

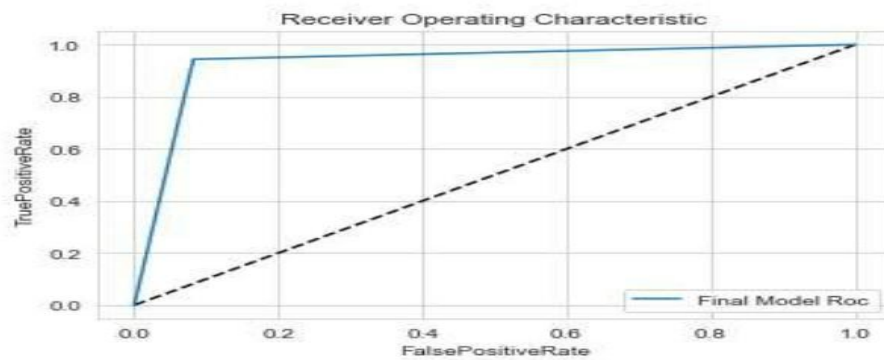


```

#Roc Curve for final model,
y_pred_fin = model.predict(x_test)
fpr , tpr, thresholds = roc_curve(y_test, y_pred_fin)

plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr1, tpr1, label= "Final Model Roc")
plt.legend()
plt.xlabel("FalsePositiveRate")
plt.ylabel("TruePositiveRate")
plt.title('Receiver Operating Characteristic')
plt.show()

```



CONCLUSION

- Key Findings and Conclusions of the Study

Online hate, described as abusive language, aggression, cyberbullying, hatefulness and many others has been identified as a major threat on online social media platforms. Social media platforms are the most prominent grounds for such toxic behaviour.

From the above analysis the below mentioned results were achieved which depicts the chances and conditions of a comment being a hateful comment or a normal comment.

- Learning Outcomes of the Study in respect of Data Science

It is possible to differentiate the comments into Malignant and Non – Malignant. However, using this project will help to create awareness among the people. It will help people to stop spreading hatred to people.

- Limitations of this work and Scope for Future Work

This project is different than the previous project provided by Flip-Robo technologies as it is text classifier using ML techniques which is challenging.

Models like decision tree classifier has taken more time and random forest and SVC algorithms are taking more time so, I didn't include those algorithms.